

# Least Squares and MLE (Gaussian Noise)

## Least Squares

### Objective Function

$$\mathcal{L}(\mathbf{w}) = \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

## MLE (Gaussian Noise)

### Likelihood

$$p(\mathbf{y} | \mathbf{X}, \mathbf{w}) = \frac{1}{(2\pi\sigma^2)^{N/2}} \prod_{i=1}^N \exp\left(-\frac{(y_i - \mathbf{w} \cdot \mathbf{x}_i)^2}{2\sigma^2}\right)$$

For estimating  $\mathbf{w}$ , the **negative log-likelihood** under Gaussian noise has the same form as the least squares objective

Alternatively, we can model the data (only  $y_i$ -s) as being generated from a distribution defined by exponentiating the negative of the objective function

## What Data Model Produces the Ridge Objective?

We have the Ridge Regression Objective, let  $\mathcal{D} = \langle (\mathbf{x}_i, y_i) \rangle_{i=1}^N$  denote the data

$$\mathcal{L}_{\text{ridge}}(\mathbf{w}; \mathcal{D}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) + \lambda \mathbf{w}^\top \mathbf{w}$$

Let's rewrite this objective slightly, scaling by  $\frac{1}{2\sigma^2}$  and setting  $\lambda = \frac{\sigma^2}{\tau^2}$ . To avoid ambiguity, we'll denote this by  $\tilde{\mathcal{L}}$

$$\tilde{\mathcal{L}}_{\text{ridge}}(\mathbf{w}; \mathcal{D}) = \frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) + \frac{1}{2\tau^2} \mathbf{w}^\top \mathbf{w}$$

Let  $\Sigma = \sigma^2 \mathbf{I}_N$  and  $\Lambda = \tau^2 \mathbf{I}_D$ , where  $\mathbf{I}_m$  denotes the  $m \times m$  identity matrix

$$\tilde{\mathcal{L}}_{\text{ridge}}(\mathbf{w}) = \frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top \Sigma^{-1} (\mathbf{y} - \mathbf{X}\mathbf{w}) + \frac{1}{2} \mathbf{w}^\top \Lambda^{-1} \mathbf{w}$$

Taking the negation of  $\tilde{\mathcal{L}}_{\text{ridge}}(\mathbf{w}; \mathcal{D})$  and exponentiating gives us a non-negative function of  $\mathbf{w}$  and  $\mathcal{D}$  which after normalisation gives a density function

$$f(\mathbf{w}; \mathcal{D}) = \exp \left( -\frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top \Sigma^{-1} (\mathbf{y} - \mathbf{X}\mathbf{w}) \right) \cdot \exp \left( -\frac{1}{2} \mathbf{w}^\top \Lambda^{-1} \mathbf{w} \right)$$

## Bayesian Linear Regression (and connections to Ridge)

Let's start with the form of the density function we had on the previous slide and factor it.

$$f(\mathbf{w}; \mathbf{D}) = \exp\left(-\frac{1}{2}(\mathbf{y} - \mathbf{X}\mathbf{w})^\top \Sigma^{-1}(\mathbf{y} - \mathbf{X}\mathbf{w})\right) \cdot \exp\left(-\frac{1}{2}\mathbf{w}^\top \Lambda^{-1}\mathbf{w}\right)$$

We'll treat  $\sigma$  as **fixed** and not treat it as a parameter. Up to a constant factor (which doesn't matter when optimising w.r.t.  $\mathbf{w}$ ), we can rewrite this as

$$\underbrace{p(\mathbf{w} \mid \mathbf{X}, \mathbf{y})}_{\text{posterior}} \propto \underbrace{\mathcal{N}(\mathbf{y} \mid \mathbf{X}\mathbf{w}, \Sigma)}_{\text{Likelihood}} \cdot \underbrace{\mathcal{N}(\mathbf{w} \mid \mathbf{0}, \Lambda)}_{\text{prior}}$$

where  $\mathcal{N}(\cdot \mid \boldsymbol{\mu}, \Sigma)$  denotes the density of the multivariate normal distribution with mean  $\boldsymbol{\mu}$  and covariance matrix  $\Sigma$

- ▶ What the ridge objective is actually finding is the **maximum a posteriori** or (MAP) estimate which is a mode of the posterior distribution
- ▶ The linear model is as described before with Gaussian noise
- ▶ The prior distribution on  $\mathbf{w}$  is assumed to be a spherical Gaussian

# Bayesian Machine Learning

In the **discriminative framework**, we model the output  $y$  as a probability distribution given the input  $\mathbf{x}$  and the parameters  $\mathbf{w}$ , say  $p(y | \mathbf{w}, \mathbf{x})$

In the Bayesian view, we assume a prior on the parameters  $\mathbf{w}$ , say  $p(\mathbf{w})$

This prior represents a “belief” about the model; the uncertainty in our knowledge is expressed mathematically as a probability distribution

When observations,  $\mathcal{D} = \langle (\mathbf{x}_i, y_i) \rangle_{i=1}^N$  are made the belief about the parameters  $\mathbf{w}$  is updated using Bayes' rule

## Bayes Rule

$$\text{For events } A, B, \Pr[A | B] = \frac{\Pr[B | A] \cdot \Pr[A]}{\Pr[B]}$$

The posterior distribution on  $\mathbf{w}$  given the data  $\mathcal{D}$  becomes:

$$p(\mathbf{w} | \mathcal{D}) \propto p(\mathbf{y} | \mathbf{w}, \mathbf{X}) \cdot p(\mathbf{w})$$

# Full Bayesian Prediction

Let us recall the posterior distribution over parameters  $\mathbf{w}$  in the Bayesian approach

$$\underbrace{p(\mathbf{w} \mid \mathbf{X}, \mathbf{y})}_{\text{posterior}} \propto \underbrace{p(\mathbf{y} \mid \mathbf{X}, \mathbf{w})}_{\text{likelihood}} \cdot \underbrace{p(\mathbf{w})}_{\text{prior}}$$

- ▶ If we use the MAP estimate, as we get more samples the posterior peaks at the MLE
- ▶ When, data is scarce rather than picking a single estimator (like MAP) we can sample from the full posterior

For  $\mathbf{x}_{\text{new}}$ , we can output the entire distribution over our prediction  $\hat{y}$  as

$$p(y \mid \mathcal{D}) = \int_{\mathbf{w}} \underbrace{p(y \mid \mathbf{w}, \mathbf{x}_{\text{new}})}_{\text{model}} \cdot \underbrace{p(\mathbf{w} \mid \mathcal{D})}_{\text{posterior}} d\mathbf{w}$$

This **integration** is often computationally very hard!

## Summary : Bayesian Machine Learning

In the Bayesian view, in addition to modelling the output  $y$  as a random variable given the parameters  $\mathbf{w}$  and input  $\mathbf{x}$ , we also encode prior belief about the parameters  $\mathbf{w}$  as a probability distribution  $p(\mathbf{w})$ .

- ▶ If the prior has a parametric form, they are called **hyperparameters**
- ▶ The posterior over the parameters  $\mathbf{w}$  is updated given data
- ▶ Either pick point (plugin) estimates, *e.g.*, **maximum a posteriori**
- ▶ Or as in the full Bayesian approach use the entire posterior to make prediction (this is often computationally intractable)
- ▶ How to choose the prior?

## How to Choose Hyper-parameters?

- ▶ So far, we were just trying to estimate the parameters  $w$
- ▶ For Ridge Regression or Lasso, we need to choose  $\lambda$
- ▶ If we perform basis expansion
  - ▶ For kernels, we need to pick the width parameter  $\gamma$
  - ▶ For polynomials, we need to pick degree  $d$
- ▶ For more complex models there may be more hyperparameters

## Using a Validation Set

- ▶ Divide the data into parts: **training**, **validation** (and **testing**)
- ▶ Grid Search: Choose values for the hyperparameters from a finite set
- ▶ Train model using **training** set and evaluate on **validation** set


$\lambda$	training error(%)	validation error(%)
0.01	0	89
0.1	0	43
1	2	12
10	10	8
100	25	27

- ▶ Pick the value of  $\lambda$  that minimises validation error
- ▶ Typically, split the data as 80% for training, 20% for validation



# Training and Validation Curves


- ▶ Plot of training and validation error vs  $\lambda$  for Lasso
- ▶ Validation error curve is  $U$ -shaped



## $K$ -Fold Cross Validation

When data is scarce, instead of splitting as training and validation:

- ▶ Divide data into  $K$  parts
- ▶ Use  $K - 1$  parts for training and 1 part as validation
- ▶ Commonly set  $K = 5$  or  $K = 10$
- ▶ When  $K = N$  (the number of datapoints), it is called LOOCV (Leave one out cross validation)



# Logistic Regression (LR)

- ▶ LR builds up on a linear model, composed with a sigmoid function


$$p(y \mid \mathbf{w}, \mathbf{x}) = \text{Bernoulli}(\text{sigmoid}(\mathbf{w} \cdot \mathbf{x}))$$

- ▶  $Z \sim \text{Bernoulli}(\theta)$

$$Z = \begin{cases} 1 & \text{with probability } \theta \\ 0 & \text{with probability } 1 - \theta \end{cases}$$

- ▶ Recall that the sigmoid function is defined by:

$$\text{sigmoid}(t) = \frac{1}{1 + e^{-t}}$$



- ▶ As we did in the case of linear models, we assume  $x_0 = 1$  for all datapoints, so we do not need to handle the bias term  $w_0$  separately

## Prediction Using Logistic Regression

Suppose we have estimated the model parameters  $\mathbf{w} \in \mathbb{R}^D$

For a new datapoint  $\mathbf{x}_{\text{new}}$ , the model gives us the probability



$$p(y_{\text{new}} = 1 \mid \mathbf{x}_{\text{new}}, \mathbf{w}) = \text{sigmoid}(\mathbf{w} \cdot \mathbf{x}_{\text{new}}) = \frac{1}{1 + \exp(-\mathbf{x}_{\text{new}} \cdot \mathbf{w})}$$

In order to make a prediction we can simply use a threshold at  $\frac{1}{2}$

$$\hat{y}_{\text{new}} = \mathbb{I}(\text{sigmoid}(\mathbf{w} \cdot \mathbf{x}_{\text{new}})) \geq \frac{1}{2}) = \mathbb{I}(\mathbf{w} \cdot \mathbf{x}_{\text{new}} \geq 0)$$

Class boundary is linear (separating hyperplane)

# Prediction Using Logistic Regression



# Likelihood of Logistic Regression

Data  $\mathcal{D} = \langle (\mathbf{x}_i, y_i) \rangle_{i=1}^N$ , where  $\mathbf{x}_i \in \mathbb{R}^D$  and  $y_i \in \{0, 1\}$

Let us denote the sigmoid function by  $\sigma$

We can write the likelihood for of observing the data given model parameters  $\mathbf{w}$  as:

$$p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = \prod_{i=1}^N \sigma(\mathbf{w}^\top \mathbf{x}_i)^{y_i} \cdot (1 - \sigma(\mathbf{w}^\top \mathbf{x}_i))^{1-y_i}$$

Let us denote  $\mu_i = \sigma(\mathbf{w}^\top \mathbf{x}_i)$

We can write the negative log-likelihood as:

$$\text{NLL}(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = - \sum_{i=1}^N (y_i \log \mu_i + (1 - y_i) \log(1 - \mu_i))$$

# Likelihood of Logistic Regression

Recall that  $\mu_i = \sigma(\mathbf{w}^\top \mathbf{x}_i)$  and the negative log-likelihood is

$$\text{NLL}(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = - \sum_{i=1}^N (y_i \log \mu_i + (1 - y_i) \log(1 - \mu_i))$$

Let us focus on a single datapoint, the contribution to the negative log-likelihood is

$$\text{NLL}(y_i \mid \mathbf{x}_i, \mathbf{w}) = -(y_i \log \mu_i + (1 - y_i) \log(1 - \mu_i))$$

This is basically the **cross-entropy** between  $y_i$  and  $\mu_i$

If  $y_i = 1$ , then as

- ▶ As  $\mu_i \rightarrow 1$ ,  $\text{NLL}(y_i \mid \mathbf{x}_i, \mathbf{w}) \rightarrow 0$
- ▶ As  $\mu_i \rightarrow 0$ ,  $\text{NLL}(y_i \mid \mathbf{x}_i, \mathbf{w}) \rightarrow \infty$

## Maximum Likelihood Estimate for LR

Recall that  $\mu_i = \sigma(\mathbf{w}^\top \mathbf{x}_i)$  and the negative log-likelihood is

$$\text{NLL}(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = - \sum_{i=1}^N (y_i \log \mu_i + (1 - y_i) \log(1 - \mu_i))$$

We can take the gradient with respect to  $\mathbf{w}$

$$\nabla_{\mathbf{w}} \text{NLL}(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) = \sum_{i=1}^N \mathbf{x}_i (\mu_i - y_i) = \mathbf{X}^\top (\boldsymbol{\mu} - \mathbf{y})$$

And the Hessian is given by,

$$\mathbf{H} = \mathbf{X}^\top \mathbf{S} \mathbf{X}$$

$\mathbf{S}$  is a diagonal matrix where  $S_{ii} = \mu_i(1 - \mu_i)$




# Calculus Background: Gradients

$$z = f(w_1, w_2) = \frac{w_1^2}{a^2} + \frac{w_2^2}{b^2}$$

$$\frac{\partial f}{\partial w_1} = \frac{2w_1}{a^2}$$

$$\frac{\partial f}{\partial w_2} = \frac{2w_2}{b^2}$$

$$\nabla_{\mathbf{w}} f = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \frac{\partial f}{\partial w_2} \end{bmatrix} = \begin{bmatrix} \frac{2w_1}{a^2} \\ \frac{2w_2}{b^2} \end{bmatrix}$$




- ▶ Gradient vectors are orthogonal to contour curves
- ▶ Gradient points in the direction of steepest increase

# Calculus Background: Hessians

$$z = f(w_1, w_2) = \frac{w_1^2}{a^2} + \frac{w_2^2}{b^2}$$

$$\nabla_{\mathbf{w}} f = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \frac{\partial f}{\partial w_2} \end{bmatrix} = \begin{bmatrix} \frac{2w_1}{a^2} \\ \frac{2w_2}{b^2} \end{bmatrix}$$


$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial w_1^2} & \frac{\partial^2 f}{\partial w_1 \partial w_2} \\ \frac{\partial^2 f}{\partial w_2 \partial w_1} & \frac{\partial^2 f}{\partial w_2^2} \end{bmatrix} = \begin{bmatrix} \frac{2}{a^2} & 0 \\ 0 & \frac{2}{b^2} \end{bmatrix}$$



- ▶ As long as all second derivatives exist, the Hessian  $H$  is symmetric
- ▶ Hessian captures the curvature of the surface

## Calculus Background: Chain Rule

$$z = f(w_1(\theta_1, \theta_2), w_2(\theta_1, \theta_2))$$



$$\frac{\partial f}{\partial \theta_1} = \frac{\partial f}{\partial w_1} \cdot \frac{\partial w_1}{\partial \theta_1} + \frac{\partial f}{\partial w_2} \cdot \frac{\partial w_2}{\partial \theta_1}$$


We will use this a lot when we study neural networks and back propagation

# General Form for Gradient and Hessian

Suppose  $\mathbf{w} \in \mathbb{R}^D$  and  $f : \mathbb{R}^D \rightarrow \mathbb{R}$

The gradient vector contains all first order partial derivatives

$$\nabla_{\mathbf{w}} f(\mathbf{w}) = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \frac{\partial f}{\partial w_2} \\ \vdots \\ \frac{\partial f}{\partial w_D} \end{bmatrix}$$



Hessian matrix of  $f$  contains all second order partial derivatives.

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial w_1^2} & \frac{\partial^2 f}{\partial w_1 \partial w_2} & \cdots & \frac{\partial^2 f}{\partial w_1 \partial w_D} \\ \frac{\partial^2 f}{\partial w_2 \partial w_1} & \frac{\partial^2 f}{\partial w_2^2} & \cdots & \frac{\partial^2 f}{\partial w_2 \partial w_D} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial w_D \partial w_1} & \frac{\partial^2 f}{\partial w_D \partial w_2} & \cdots & \frac{\partial^2 f}{\partial w_D^2} \end{bmatrix}$$

# Gradient Descent Algorithm



**Gradient descent** is one of the simplest, but very general algorithm for optimization

It is an iterative algorithm, producing a new  $\mathbf{w}_{t+1}$  at each iteration as

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \mathbf{g}_t = \mathbf{w}_t - \eta_t \nabla f(\mathbf{w}_t)$$



We will denote the gradients by  $\mathbf{g}_t$

$\eta_t > 0$  is the **learning rate** or **step size**



## Choosing a Step Size

- ▶ Choosing a good step-size is important
- ▶ If step size is too large, algorithm may never converge
- ▶ If step size is too small, convergence may be very slow
- ▶ May want a time-varying step size



## Iteratively Re-Weighted Least Squares (IRLS)

Depending on the dimension, we can apply Newton's method to estimate  $\mathbf{w}$

Let  $\mathbf{w}_t$  be the parameters after  $t$  Newton steps.

The gradient and Hessian are given by:

$$\begin{aligned}\mathbf{g}_t &= \mathbf{X}^\top(\boldsymbol{\mu}_t - \mathbf{y}) = -\mathbf{X}^\top(\mathbf{y} - \boldsymbol{\mu}_t) \\ \mathbf{H}_t &= \mathbf{X}^\top \mathbf{S}_t \mathbf{X}\end{aligned}$$

The Newton Update Rule is:

$$\begin{aligned}\mathbf{w}_{t+1} &= \mathbf{w}_t - \mathbf{H}_t^{-1} \mathbf{g}_t \\ &= \mathbf{w}_t + (\mathbf{X}^\top \mathbf{S}_t \mathbf{X})^{-1} \mathbf{X}^\top (\mathbf{y} - \boldsymbol{\mu}_t) \\ &= (\mathbf{X}^\top \mathbf{S}_t \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{S}_t (\mathbf{X} \mathbf{w}_t + \mathbf{S}_t^{-1} (\mathbf{y} - \boldsymbol{\mu}_t)) \\ &= (\mathbf{X}^\top \mathbf{S}_t \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{S}_t \mathbf{z}_t\end{aligned}$$

Where  $\mathbf{z}_t = \mathbf{X} \mathbf{w}_t + \mathbf{S}_t^{-1} (\mathbf{y} - \boldsymbol{\mu}_t)$ . Then  $\mathbf{w}_{t+1}$  is a solution of the following:

### Weighted Least Squares Problem

$$\text{minimise } \sum_{i=1}^N S_{t,ii} (z_{t,i} - \mathbf{w}^\top \mathbf{x}_i)^2$$

# Multiclass Logistic Regression

Multiclass logistic regression is also a discriminative classifier

Let the inputs be  $\mathbf{x} \in \mathbb{R}^D$  and  $y \in \{1, \dots, C\}$

There are parameters  $\mathbf{w}_c \in \mathbb{R}^D$  for every class  $c = 1, \dots, C$

We'll put this together in a matrix form  $\mathbf{W}$  that is  $D \times C$

The multiclass logistic model is given by:

$$p(y = c \mid \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_c^\top \mathbf{x})}{\sum_{c'=1}^C \exp(\mathbf{w}_{c'}^\top \mathbf{x})}$$



## Multiclass Logistic Regression

The multiclass logistic model is given by:

$$p(y = c \mid \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_c^T \mathbf{x})}{\sum_{c'=1}^C \exp(\mathbf{w}_{c'}^T \mathbf{x})}$$

Recall the softmax function

### Softmax

Softmax maps a set of numbers to a probability distribution with mode at the maximum

$$\text{softmax} \left( [a_1, \dots, a_C]^T \right) = \left[ \frac{e^{a_1}}{Z}, \dots, \frac{e^{a_C}}{Z} \right]^T$$

$$\text{where } Z = \sum_{c=1}^C e^{a_c}.$$

The multiclass logistic model is simply:

$$p(y \mid \mathbf{x}, \mathbf{W}) = \text{softmax} \left( \left[ \mathbf{w}_1^T \mathbf{x}, \dots, \mathbf{w}_C^T \mathbf{x} \right]^T \right)$$

# Multiclass Logistic Regression

